

Non-Convex Polygon Packing via Simulated Annealing and Bracketing-Bisection Search

Joel Maldonado*

April 2026

Abstract. The problem, featured in a 2024 Kaggle competition, is to find the minimum side length L^* of a square container that can hold N congruent copies without overlap, each free to translate and rotate. Answering this exactly is NP-hard. We describe a practical solver combining time-budgeted embarrassingly parallel Simulated Annealing with an outer bisection loop that narrows the bracket containing L^* , followed by a feasibility-polish stage. Collision detection uses the Separating Axis Theorem (SAT), chosen because it returns penetration depth, converting a combinatorial feasibility question into a smooth, navigable energy function. We deployed embarrassingly parallel jobs on a node of the HPC cluster for $N \in \{5, 10, 25, 50, 100, 200\}$, with wall times from 3 minutes to 8 hours. Packing densities range from $\eta \approx 0.58$ at small N to $\eta \approx 0.48$ at $N = 200$.

Keywords. polygon packing; simulated annealing; separating axis theorem; bisection search; embarrassingly parallel HPC; packing density scaling.

1. INTRODUCTION

Fix a non-convex polygon, a 15-vertex shape and make N identical copies. The problem is to find the smallest square that holds all of them without overlap. As N grows, the configuration space expands exponentially and exact enumeration becomes hopeless. The simpler case of unit disks packing decision problem is NP-complete. If the decision version of a problem is NP-complete, then its optimization version is NP-hard. Non-convexity adds a further layer of difficulty. It makes collision detection more expensive, the energy landscape more rugged and potentially fragmented into disconnected islands of feasible configurations. Exact geometric approaches exist. The No-Fit Polygon (NFP), surveyed by Bennell and Oliveira [1], provides a precise characterization of relative placements between shapes. However, for non-convex polygons it typically requires decomposition and non-trivial geometric machinery. The level of implementation complexity of the NFP approach is a risk to the fast prototyping needed for a competition. Thus we opt for a method both simpler to implement and flexible but with the drawback of providing only approximate solutions.

We therefore chose to implement a penalty-based Simulated Annealing (SA) approach. Instead of restricting the search to strictly feasible configurations, we allow intermediate overlap and penalize it continuously. This relaxes the geometry of the search space: rather than navigating only through narrow feasible corridors, the solver can move through infeasible regions while still being guided toward feasibility.

Other metaheuristics provide natural alternatives. Hybrid SA-genetic approaches have been applied to orthogonal packing problems [5], and population-based methods are widely used in cutting and packing to explore multiple candidate layouts in parallel. Parallel tempering is also particularly appealing for rugged landscapes: it runs replicas at different temperatures and exchanges configurations via a Metropolis criterion, allowing the system to mix more effectively across energy barriers.

For this quick prototyping study aimed at a competition, simulated annealing by itself is the right starting point. As Sokal famously put it: “Monte Carlo is an extremely bad method; it should be used only when all alternative methods are worse.” What that observation misses is that its very simplicity makes it a reliable fallback that works across a broad class of problems and is often a good reliable baseline. For us, simulated annealing is not a last resort so much as a deliberate foundation: it maps cleanly onto GPU architecture, its energy landscape connects naturally to statistical mechanics and the Ising model, and understanding its behavior on this geometry is the necessary groundwork before introducing the added complexity of population-based or replica methods.

*Department of Mathematics, University of Arizona. joelmaldonado@arizona.edu. Supported by allocation math589 on the University of Arizona HPC cluster.

2. PROBLEM FORMULATION

Let $\mathcal{P}$ be a fixed non-convex polygon with 15 vertices and area $A_p = 0.245625$, computed by the shoelace formula from the reference vertex array. Each copy i is placed by specifying its centre (x_i, y_i) and rotation angle $\theta_i \in [0, 2\pi)$; world vertices follow as

$$w_k^{(i)} = R(\theta_i) v_k + c_i.$$

A configuration is *feasible* if no two copies overlap and every copy lies within $\Omega_L = [-L/2, L/2]^2$. The goal is to find, for each N , the smallest L admitting a feasible configuration.

Since all N copies must fit inside Ω_L without overlap, the square area satisfies $L^2 \geq NA_p$, giving the hard lower bound

$$L_{\text{area}} = \sqrt{NA_p}.$$

The bisection loop never probes below this. Our Simulated Annealing implementation performs a move by selecting a random copy i and proposing a new state (x'_i, y'_i, θ'_i) by adding a random offset to the current state, with step sizes (s_{xy}, s_θ) that adapt over time to maintain a target acceptance rate. The energy change ΔE is computed incrementally by checking only pairs involving copy i , and the move is accepted with probability $\min(1, e^{-\Delta E/T})$ where T is the current temperature. If the energy is within the feasibility threshold, the configuration is saved and the bisection loop updates its bracket on L^* accordingly. After bisection, an adaptive polish phase continues to attempt smaller L values until the wall time budget is exhausted, with the best known L^* updated on every improvement. Restricting the search to strictly non-overlapping configurations at all times confines it to disconnected feasible islands in $\mathbb{R}^{2N} \times [0, 2\pi)^N$. To overcome this problem we allow a budget for some overlap, penalizing it continuously in the energy function. This then guides our choice of a continuous collision metric instead of a binary collision test.

3. ALGORITHM

3.1. Collision Detection Pipeline

During our Simulated Annealing pipeline, collision detection is invoked after every proposed move to evaluate the energy change, making it the dominant runtime cost. Three collision detection mechanisms are employed: bounding circle checks, axis-aligned bounding box (AABB) checks, and the Separating Axis Theorem (SAT). In this order each detection becomes more expensive. Having a mixture of cheap and expensive checks allows us to filter most detections early, so that the expensive SAT is only invoked on a small fraction of pairs that pass the broad-phase filters. Naive pairwise checking costs $O(N^2)$ per iteration. A uniform spatial hash grid with cell size double the size of the bounding radius

$$h = 2r_{\text{br}} = 1.6$$

restricts each query to a 5×5 neighbourhood of $\approx 261\eta$ candidates, where

$$\eta = NA_p/L^2$$

is the packing density. The energy change is then computed incrementally, touching only pairs involving the moved polygon, reducing the effective per-move cost to $O(1)$. Of those candidates, A large fraction of these candidates is expected to be discarded by the polygon-level AABB checks. Group-level and per-triangle AABB checks prune the remainder before the Separating Axis Theorem (SAT) is invoked. Figure 1 visualizes each phase of the four-stage collision detections pipeline.

Offline, the polygon is triangulated as a fan from vertex 0 into $T = 13$ triangles. These are grouped into three clusters of sizes $(5, 5, 3)$ and cached as AABBs. During narrow-phase evaluation, SAT projects triangle pairs onto six edge-normal axes to compute the minimum overlap width, $d \geq 0$. World-space vertices are projected directly to avoid recomputing rotation matrices on each call. If a gap appears on any axis, the triangles are disjoint. Otherwise, $d > 0$ yields the penetration depth, providing a continuous measure of intersection severity. Relying on a binary collision penalty would create a flat energy landscape with no gradient. By contrast, the depth d supplies a clear directional signal. Then, feeding d^2 into the energy function generates the smooth guidance necessary to steer the simulated annealing optimization toward feasibility from arbitrary starting states.

3.2. Energy Function and Simulated Annealing Schedule

We model feasibility as a soft-constrained optimization problem, penalizing both geometric interpenetration and container boundary violations. The total energy of a configuration at container side length L is

$$E = \lambda \sum_{i < j} \sum_{(a,b) \in \mathcal{T}_{ij}} d(a,b)^2 + \mu \sum_i \phi_i(L),$$

Collision Detection Pipeline




Figure 1. Collision detection pipeline. Top-left: broad phase using a uniform spatial hash and neighborhood query. Top-right: polygon AABB culling. Bottom-left: group-level AABB pruning after triangulation into clustered subsets. Bottom-right: triangle-triangle SAT narrow phase returning penetration depth d , which feeds the overlap energy.

Table 1. Two-phase simulated annealing schedule parameters.

Parameter	Phase A	Phase B
Iterations M	70,000	35,000
T_{start}	0.30	0.06
T_{end}	3×10^{-4}	10^{-6}
Overlap power p	1	2
λ_0, μ_0	10^2 (fixed)	10^3 (ramped)
Step adapt window	1,500 iters	
Acceptance band	[0.25, 0.55]	[0.10, 0.30]
Step shrink / grow	0.88 / 1.12	0.90 / 1.08
Reinsertion prob.	2%	0.5%
Feasibility Threshold $\mathcal{F}$	10^{-6}	

where $d(a, b)$ is the SAT penetration depth between triangle pair (a, b) , $\phi_i(L)$ is the squared AABB boundary violation of copy i , and (λ, μ) are the respective penalty weights. A configuration is declared strictly feasible when the unweighted residual is negligible:

$$\mathcal{F} = \sum_{i < j} \text{overlap}_{ij} + \sum_i \phi_i < 10^{-6}.$$

To minimize E , the SA trial executes two sequential phases with geometrically decreasing temperature

$$T_k = T_0 \beta^k, \text{ where } \beta = (T_{\text{end}}/T_0)^{1/M}.$$

Phase A (Exploration; 70,000 iterations): Prioritizes broad search using a high initial temperature $T_0 = 0.30$ and a linear overlap penalty ($p = 1$). Phase B (Feasibility; 35,000 iterations): Drives the system toward strict feasibility by dropping T_0 to 0.06, switching to a quadratic penalty ($p = 2$), and scaling penalty weights by a factor of 1.8 every 1,200 iterations. Throughout both phases, step sizes (s_{xy}, s_θ) dynamically adapt every 1,500 iterations to maintain a target acceptance band, automatically coupling the search radius to the cooling temperature. While small N instances execute the full 105,000-iteration budget, wall-clock limits may truncate Phase B for larger N . Table 1 summarizes the schedule.

3.3. Outer Loop: Bisection and Polish

The Simulated Annealing solver serves as a feasibility oracle: for a fixed container side length L , it answers whether a strictly feasible packing can be found. Because simulated annealing does not directly minimize L , we wrap it in an outer bisection loop.

In the 2025 implementation, bracketing does not start from a fixed $[L_{\text{area}}, \gamma L_{\text{area}}]$ interval. Instead, the run starts from a grid-based initialization and performs shrink/grow passes to produce a valid bracket with L_{low} infeasible and L_{high} feasible. Values at or below $L_{\text{area_infeas}} = \sqrt{NA_p}(1 - 10^{-9})$ are treated as provably infeasible.

After bracketing, bisection is run for a fixed 26 iterations. At each probe, $L_{\text{mid}} = \frac{1}{2}(L_{\text{low}} + L_{\text{high}})$ is tested with SA; feasible updates $L_{\text{high}} \leftarrow L_{\text{mid}}$, infeasible updates $L_{\text{low}} \leftarrow L_{\text{mid}}$.

To reduce restart cost between consecutive probes, we use warm starts. When a previous feasible state exists at L_{old} , polygon centers are scaled by $\alpha = L_{\text{new}}/L_{\text{old}}$ before the next simulated annealing run, preserving layout structure and reducing the time needed to recover feasibility.

Bisection + warm-start pseudocode

1. Build bracket by shrink/grow passes from a grid-initialized state: find infeasible L_{low} and feasible L_{high} .
2. For $k = 1, \dots, 26$:
 - (a) $L_{\text{mid}} \leftarrow \frac{1}{2}(L_{\text{low}} + L_{\text{high}})$.
 - (b) If $L_{\text{mid}} \leq L_{\text{area_infeas}}$, set $L_{\text{low}} \leftarrow L_{\text{mid}}$ and continue.
 - (c) Warm-start by scaling centers from previous feasible L to L_{mid} .
 - (d) Run simulated annealing feasibility probe at L_{mid} .
 - (e) If feasible: $L_{\text{high}} \leftarrow L_{\text{mid}}$, save warm state, update best feasible L^* .
 - (f) Else: $L_{\text{low}} \leftarrow L_{\text{mid}}$.

Polish-stage pseudocode

1. Start from post-bisection best feasible state $(L^*, \mathcal{C}^*)$.
2. Maintain adaptive shrink factor ε and a rolling window of 20 attempts.
3. Repeatedly propose $L_{\text{try}} = L^*(1 - \varepsilon)$, warm-start, and run SA.
4. On success: accept L_{try} , update $(L^*, \mathcal{C}^*)$.

Table 2. Best feasible side length, packing efficiency, and feasibility rate by instance size. $N = 5$ and $N = 10$ used the five-node pipeline (3-min smoke tests); $N \geq 25$ used single-node array jobs (4–8 h). Cross- N comparisons should be read with this provenance difference in mind.

N	L_{area}	L^*	η	Runs	t_{first}
5	1.108	1.472	0.567	15	5 s
10	1.568	2.056	0.581	15	6 s
25	2.478	3.315	0.559	6	8 s
50	3.504	4.652	0.567	4	46 s
100	4.956	6.785	0.533	3	53 s
200	7.009	10.104	0.481	1	19 s

5. Every 20 attempts, compare success rate to the 35% target and adapt ϵ .
6. Flush checkpoints periodically and on SIGTERM.

After bisection, an adaptive polish phase progressively shrinks the container: at each attempt the solver tries $L \leftarrow L^*(1 - \epsilon)$ with ϵ adapted over rolling windows of 20 probes to maintain a target success rate of roughly 35%, then runs simulated annealing at the candidate L . A successful packing updates the best known L^* . Artifacts confirm this phase produces repeated improvements in L over time, and it accounts for the majority of density gain at large N . Progress is written to disk on every improvement, and a global snapshot is flushed on SIGTERM so SLURM preemption loses no work. Every worker is seeded deterministically from (base seed, run ID, trial index) via SplitMix64, guaranteeing full reproducibility.

4. EXPERIMENTAL SETUP

All computational experiments were conducted on the University of Arizona high-performance computing (HPC) cluster. The results reported here were executed as single-node SLURM array jobs. Each array task executed a single independent random seed, with a maximum of two seeds running concurrently per job. We configured the array sizes and wall-clock limits according to the problem size: 8 tasks at $N = 25$ (4-hour wall time, yielding an effective search budget of 14,100,s after a 300-second post-processing cushion), 6 tasks at $N = 50$ (4,h), 4 tasks at $N = 100$ (4,h), and 2 tasks at $N = 200$ (8,h wall time, yielding a 28,500,s effective budget). To ensure data integrity, SLURM was configured to send a SIGTERM signal 120,s prior to the hard time limit, allowing the solver sufficient time to flush its best-found snapshot to disk. Because these runs executed without inter-worker communication, the global best solution for each instance of N was identified through a post-hoc reduction pass over all resulting output CSV files.

5. RESULTS

We first inspect the best feasible side-length trajectories over wall time, then the final recovered packing configurations and their packing densities, and finally a summary of the best feasible side lengths.

Table 2 reports the best feasible side length L^* and packing efficiency $\eta = NA_p/(L^*)^2$ for each N . The area lower bound $L_{\text{area}} = \sqrt{NA_p}$ is given for reference; $\eta = 1$ would require zero wasted space, which is geometrically impossible for this non-convex shape. The column “Runs” counts workers that found at least one feasible configuration. It is possible that no feasible configuration was found within the time budget. This is the reason we run multiple independent trials per N : to increase the probability of finding at least one feasible configuration. If we find multiple feasible configurations across different runs, we report the best one.

Observe three patterns emerge. First, the number of workers reaching feasibility drops sharply: 15 of 15 at $N = 5$ and $N = 10$, down to 6 at $N = 25$, and to a single run at $N = 200$. The simulated annealing is time-budgeted and not iteration-bounded, but the cost per simulated annealing proposal grows with N since more polygon pairs must be checked. At large N we can fit fewer trials within the wall-time budget, and the trials that do run explore a configuration space of dimension $3N$ with proportionally less coverage per polygon. The configuration space at $N = 200$ has 600 degrees of freedom.

Second, the best feasible side length increases monotonically with N , while the gap between L^* and the area lower bound remains substantial across all tested sizes. The final packings in Figure 3 show this increasingly constrained geometry directly: larger instances still admit feasible layouts, but with visibly tighter global packing structure and less slack for local rearrangement.

Third, the median time to final best configuration was 205 s at $N = 25$ but exceeded 3,250 s at $N = 100$, out of a 14,100 s budget. The trajectories show rapid early gains followed by extended plateaus with smaller late improvements. This indicates diminishing returns in long single runs and supports allocating more budget to many independent shorter trials, especially at larger N .










Figure 2. Best-so-far feasible square side length L versus wall time for $N \in \{5, 10, 25\}$.

Final Best Configurations by Instance Size

 $N = 5$  $N = 10$  $N = 25$  $N = 50$  $N = 100$  $N = 200$ **Figure 3.** Best recovered polygon packing configurations for each reported problem size.




Figure 4. Best-so-far feasible square side length L versus wall time for $N \in \{50, 100, 200\}$.

6. CONCLUSION

The interesting aspect of our approach was the use of the Separating Axis Theorem (SAT) to measure penetration depth rather than relying on a binary collision test. This allowed us to craft an energy function that allows for states with some overlap and a continuous penalty that guides the Simulated Annealing algorithm out of messy, heavily overlapping starting states and into strictly feasible packings.

We observe that many of the final gains in a configuration are achieved during the adaptive polish phase. Once the outer bisection loop rapidly narrowed down the target container size L^* , the adaptive polish phase did the heavy lifting, squeezing out the majority of the final density gains.

However, our HPC experiments revealed a fundamental bottleneck. As the number of polygons N increases, the computational cost of collision checking grows so large that our time-budgeted runs could no longer adequately explore the high-dimensional search space, causing the success rate to collapse.

The primary contribution of this study is therefore strategic: in a geometric landscape this rugged and fragmented, trial *diversity* matters vastly more than trial *duration*. Letting a single simulation run for hours will yield rapidly diminishing returns the moment it gets trapped in a deep local minimum. Although we have some wiggle room for overlap penalties, in packings with large N our budget might not be enough to find a better rearrangement, so an adaptive update rule for the penalty budget is needed. A much smarter and more effective way to reduce this bottleneck is to launch many short, independent screening runs, dedicating continued compute time only to the most promising states. This is a natural fit for GPU architecture, which excels at running large batches of independent trials in parallel. By massively scaling the number of parallel trials, we can maintain a high level of structural diversity in the search process, increasing the chances of finding high-quality packings even in the face of an exponentially growing configuration space.

Future work must utilize this lesson. First, estimate the average time needed to obtain a sound packing, then within that budget run as many independent trials as possible and select the best ones for continued refinement. Second, integrating population-based methods such as parallel tempering or genetic recombination would organically maintain the structural diversity that serial Simulated Annealing lacks. Finally, while porting this solver to a GPU architecture remains highly valuable, its true power will lie in massively scaling the volume of independent parallel trials.

REFERENCES

- [1] J. A. Bennell and J. F. Oliveira. The geometry of nesting problems: A tutorial. *Eur. J. Oper. Res.*, 184(2):397–415, 2008.
- [2] E. K. Burke, G. Kendall, and G. Whitwell. A new placement heuristic for the orthogonal stock-cutting problem. *Oper. Res.*, 52(4):655–671, 2004.
- [3] Kaggle. Santa 2025 competition: Christmas Tree Packing Challenge. *kaggle.com/competitions/santa-2025*, 2025.
- [4] S. Kirkpatrick, C. D. Gelatt, and M. P. Vecchi. Optimization by simulated annealing. *Science*, 220(4598):671–680, 1983.
- [5] T. W. Leung, C. K. Chan, and M. D. Troutt. Application of a mixed simulated annealing-genetic algorithm heuristic for the two-dimensional orthogonal packing problem. *Eur. J. Oper. Res.*, 145(3):530–542, 2003.
- [6] A. Sokal. Monte Carlo methods in statistical mechanics: Foundations and new algorithms. In *Functional Integration*, pages 131–192. Springer, 1997.
- [7] T. Vogel and Y. W. Li. Accelerating multicanonical sampling with irreversibility. *arXiv preprint arXiv:2601.18980*, 2026.
- [8] L. Zhang, P. Potapchik, J. He, Y. Du, A. Doucet, F. Vargas, H.-D. Dau, and S. Syed. Accelerated parallel tempering via neural transports. *arXiv preprint arXiv:2502.10328*, 2025.

A. APPENDIX: DENSITY TRAJECTORY




Figure 5. Best-so-far packing density $\eta = NA_p/L^2$ versus wall time. The late-time slope at larger N is consistent with most of the final improvement occurring during feasibility polish rather than during the initial bracket search.