

# **Discrete and Continuous Methods for Packing Problems**

Your Name

January 22, 2026

# Contents

# Preface

This book develops a unified framework for packing problems, combining discrete optimization, graph theory, and continuous methods from applied mathematics.

# Chapter 1

## Simulated Annealing for Circle Packing in a Square

### Guiding Question

How can a stochastic, temperature–driven algorithm reliably solve deterministic, highly non–convex geometric optimization problems where gradient–based methods fail?

This chapter establishes the *theoretical and conceptual foundation* for simulated annealing (SA) through an inquiry–based lens. We use the geometric problem of packing  $N$  identical circles into a square of minimal side length as a concrete running example. The goal is not to implement SA yet, but to understand *why it works, what assumptions it relies on, and which design choices matter* before turning to test–driven C implementations in later chapters.

### 1.1 The Optimization Problem

We consider the problem of packing  $N$  circles of radius  $r$  into a square of side length  $L$ , minimizing  $L$  subject to non–overlap and boundary constraints.

#### 1.1.1 Configuration Space

A configuration is described by

$$x = (x_1, y_1, \dots, x_N, y_N) \in \mathbb{R}^{2N},$$

with constraints

$$|(x_i, y_i) - (x_j, y_j)| \geq 2r, \quad i \neq j, \quad r \leq x_i, y_i \leq L - r.$$

The feasible set is a subset of  $\mathbb{R}^{2N}$  defined by pairwise distance constraints and box constraints.

## 1.2 Why Classical Optimization Fails

### 1.2.1 Inquiry

- Is the feasible set convex?
- Is the objective function differentiable everywhere?
- How does the number of local minima scale with  $N$ ?

Each question reveals a structural obstruction:

- The feasible set is *highly non-convex*.
- Contact events between circles introduce *nonsmoothness*.
- The number of metastable configurations grows combinatorially with  $N$ .

These are not numerical pathologies but geometric ones. Gradient-based methods fail *by design*, not merely by poor tuning.

## 1.3 From Optimization to Sampling

The key conceptual shift of simulated annealing is to replace deterministic descent with *probabilistic exploration*.

### 1.3.1 Energy Formulation

We introduce a scalar energy

$$E(x, L) = E_{\text{overlap}}(x) + \alpha L,$$

where  $E_{\text{overlap}}$  penalizes violations of geometric constraints and  $\alpha > 0$  balances feasibility against compactness.

### 1.3.2 Key Question

Why should we ever accept a move that increases  $E$ ?

The answer is central: rejecting all uphill moves traps the algorithm in local minima. Allowing energy-increasing moves enables escape from metastable states.

## 1.4 Statistical Mechanics Perspective

## Chapter 2

# Simulated Annealing for Circle Packing in a Square

### Guiding Question

How can a stochastic, temperature–driven algorithm reliably solve deterministic, highly non–convex geometric optimization problems where gradient–based methods fail?

This chapter establishes the *theoretical and conceptual foundation* for simulated annealing (SA) through an inquiry–based lens. We use the geometric problem of packing  $N$  identical circles into a square of minimal side length as a concrete running example. The goal is not to implement SA yet, but to understand *why it works, what assumptions it relies on, and which design choices matter* before turning to test–driven C implementations in later chapters.

### 2.1 The Optimization Problem

We consider the problem of packing  $N$  circles of radius  $r$  into a square of side length  $L$ , minimizing  $L$  subject to non–overlap and boundary constraints.

#### 2.1.1 Configuration Space

A configuration is described by

$$x = (x_1, y_1, \dots, x_N, y_N) \in \mathbb{R}^{2N},$$

with constraints

$$|(x_i, y_i) - (x_j, y_j)| \geq 2r, \quad i \neq j, \quad r \leq x_i, y_i \leq L - r.$$

The feasible set is a subset of  $\mathbb{R}^{2N}$  defined by pairwise distance constraints and box constraints.

## 2.2 Why Classical Optimization Fails

### 2.2.1 Inquiry

- Is the feasible set convex?
- Is the objective function differentiable everywhere?
- How does the number of local minima scale with  $N$ ?

Each question reveals a structural obstruction:

- The feasible set is *highly non-convex*.
- Contact events between circles introduce *nonsmoothness*.
- The number of metastable configurations grows combinatorially with  $N$ .

These are not numerical pathologies but geometric ones. Gradient-based methods fail *by design*, not merely by poor tuning.

## 2.3 From Optimization to Sampling

The key conceptual shift of simulated annealing is to replace deterministic descent with *probabilistic exploration*.

### 2.3.1 Energy Formulation

We introduce a scalar energy

$$E(x, L) = E_{\text{overlap}}(x) + \alpha L,$$

where  $E_{\text{overlap}}$  penalizes violations of geometric constraints and  $\alpha > 0$  balances feasibility against compactness.

### 2.3.2 Key Question

Why should we ever accept a move that increases  $E$ ?

The answer is central: rejecting all uphill moves traps the algorithm in local minima. Allowing energy-increasing moves enables escape from metastable states.

## 2.4 Statistical Mechanics Perspective

Simulated annealing is grounded in equilibrium statistical mechanics.

### 2.4.1 Gibbs Measure

For temperature  $T > 0$ , define

$$\pi_T(x) = Z_T^{-1} e^{-E(x)/T},$$

where  $Z_T$  is the normalizing constant.

- High  $T$ : broad exploration of configuration space.
- Low  $T$ : concentration near global minima.

Optimization emerges as the zero-temperature limit of sampling.

### 2.4.2 Metropolis Acceptance Rule

Given a current state  $x$  and a proposal  $x'$ , accept  $x'$  with probability

$$p = \min(1, e^{-(E(x') - E(x))/T}).$$

### 2.4.3 Inquiry

- Why does this rule preserve  $\pi_T$  as an invariant distribution?
- What role does detailed balance play?

These questions will later justify the correctness of the algorithm independent of implementation.

## 2.5 Simulated Annealing as a Limit Process

Simulated annealing proceeds by lowering the temperature:

$$T_0 > T_1 > \dots > T_k \rightarrow 0.$$

### 2.5.1 Theoretical Guarantee

With logarithmic cooling,  $T_k \sim c/\log k$ , simulated annealing converges almost surely to a global minimizer. Although impractical, this result explains *why* the method can work at all.

## 2.6 Algorithmic Ingredients

Every simulated annealing algorithm consists of:

1. **State space:** configurations  $x$  (and possibly  $L$ ).

2. **Energy:** objective plus penalties.
3. **Proposal kernel:** how candidate states are generated.
4. **Acceptance rule:** Metropolis criterion.
5. **Cooling schedule:** temperature decay.

Failure of any component compromises the method.

## 2.7 Energy Design for Circle Packing

Hard constraints destroy ergodicity. Instead we use soft penalties.

### 2.7.1 Pairwise Overlap Penalty

For distance  $d$  between circle centers,

$$\phi(d) = \begin{cases} (2r - d)^2, & d < 2r, \\ 0, & d \geq 2r. \end{cases}$$

Total overlap energy:

$$E_{\text{overlap}}(x) = \sum_{i < j} \phi(|x_i - x_j|) + \sum_i \phi_{\text{wall}}(x_i).$$

### 2.7.2 Inquiry

- Why quadratic penalties instead of infinite barriers?
- How does penalty scaling interact with temperature?

## 2.8 Proposal Mechanisms

Common proposal strategies include:

- local Gaussian perturbations of a single circle,
- occasional global reshuffling,
- temperature-dependent step sizes.

A key heuristic is

$$\text{proposal scale} \propto \sqrt{T}.$$

## 2.9 Optimizing the Square Size

Two approaches are common:

### 2.9.1 Feasibility Search

Fix  $L$ , test feasibility via SA, and perform an outer binary search on  $L$ .

### 2.9.2 Joint Optimization

Optimize  $(x, L)$  jointly using

$$E(x, L) = E_{\text{overlap}}(x) + \alpha L.$$

### 2.9.3 Inquiry

- Why must  $\alpha$  be temperature-aware?
- What failure modes arise if  $L$  shrinks too early?

## 2.10 Common Failure Modes

Symptom	Cause
Freezing	Cooling too fast
Jittering	Proposal scale too large
Poor minima	Bad initialization
Slow convergence	Poor energy scaling

## 2.11 Exercises

1. Prove detailed balance for the Metropolis acceptance rule.
2. Show that hard constraints break ergodicity.
3. Compare simulated annealing and basin hopping on small  $N$ .
4. Design a temperature-dependent penalty function.
5. Explain why SA only superficially resembles stochastic gradient descent.

## 2.12 Conceptual Summary

Simulated annealing replaces deterministic descent with controlled stochastic exploration. Its validity comes from statistical mechanics; its effectiveness comes from careful algorithmic design.

In later chapters, this conceptual foundation will be translated into *test-driven C implementations*, where correctness precedes performance.

# Chapter 3

## Test-Driven Simulated Annealing: From Theory to a Na"ive C Baseline

### Purpose

This chapter translates the conceptual foundations of simulated annealing developed in Chapter 1 into a *correctness-first implementation workflow*. The emphasis is not speed or sophistication, but **trustworthiness**: every algorithmic component is introduced via inquiry and then validated through test-driven development (TDD).

By the end of the chapter, we obtain a minimal simulated annealing solver for 2D circle packing that is:

- reproducible,
- instrumented,
- mathematically faithful to the theory,
- and suitable as a baseline for later optimization (OpenMP, CUDA, etc.).

### 3.1 How This Chapter Fits the Narrative

Chapter 1 answered the question *why* simulated annealing is appropriate for nonconvex geometric optimization. This chapter answers a different question:

*How do we implement simulated annealing so that each theoretical assumption is explicitly checked and enforced by tests?*

The organizing principle is:

*extTheory* → Algorithmic Invariant → Unit Test → Code.

## 3.2 Learning Outcomes

After completing this chapter, you should be able to:

- implement a minimal Metropolis kernel and annealing loop;
- encode geometric constraints using smooth penalty energies;
- write tests that enforce probabilistic and geometric invariants;
- diagnose failures using acceptance-rate and energy traces;
- extend the baseline toward feasibility search and joint  $(x, L)$  optimization.

## 3.3 Project Skeleton

We use a deliberately small and explicit structure. Each module corresponds to one theoretical component from Chapter 1.

```
sa_circlemesh/
include/           (C headers in later chapters)
src/
energy.py          # E(x,L)
propose.py         # proposal kernel
metropolis.py     # acceptance rule
anneal.py          # annealing loop
rng.py             # deterministic randomness
tests/
test_energy.py
test_propose.py
test_metropolis.py
test_anneal.py
test_reproducibility.py
```

Although this chapter uses Python for rapid iteration, the structure mirrors the C implementations developed later.

## 3.4 Design Specification: Naive Baseline

### 3.4.1 State

A state consists of

$$(X, L), \quad X \in \mathbb{R}^{N \times 2}, \quad L > 0,$$

where  $X$  stores circle centers and  $L$  is the square side length.

### 3.4.2 Energy

We implement exactly the energy introduced in Chapter 1:

$$E(X, L) = E_{\text{pair}}(X) + E_{\text{wall}}(X, L) + \alpha L.$$

Pairwise overlap penalty:

$$E_{\text{pair}}(X) = \sum_{i < j} \phi(|x_i - x_j|), \quad \phi(d) = \max(0, 2r - d)^2.$$

Wall penalty:

$$E_{\text{wall}}(X, L) = \sum_i \sum_{k=1}^2 \left[ \max(0, r - x_{i,k})^2 + \max(0, x_{i,k} - (L - r))^2 \right].$$

### 3.4.3 Inquiry

- Why must  $E \geq 0$  always?
- Why must feasible configurations yield  $E_{\text{pair}} = E_{\text{wall}} = 0$ ?

These questions become explicit unit tests.

### 3.4.4 Proposal Kernel

A proposal perturbs exactly one circle:

$$x'_k = x_k + \delta, \quad \delta \sim \mathcal{N}(0, \sigma^2 I_2).$$

This choice enforces locality and ergodicity. In the naive baseline, proposals may be clipped to remain within a bounding box to avoid numerical blow-up.

### 3.4.5 Acceptance Rule

Given  $\Delta E = E' - E$  at temperature  $T$ , we apply the Metropolis criterion:

$$\mathbb{P}(\text{accept}) = \min(1, e^{-\Delta E / T}).$$

### 3.4.6 Cooling Schedule

We use geometric cooling:

$$T_{k+1} = \gamma T_k, \quad \gamma \in (0, 1).$$

This is not theoretically optimal, but it is sufficient for a correctness baseline.

## 3.5 Inquiry-Driven TDD Sequence

Each exercise introduces a single invariant implied by the theory and enforces it via tests.

### 3.5.1 Exercise 1: Deterministic Randomness

**Invariant:** identical seeds imply identical trajectories.

**Test:**

- Repeated calls with the same seed produce identical proposals and acceptance decisions.

### 3.5.2 Exercise 2: Energy Sanity

**Invariants:**

1.  $E(X, L) \geq 0$  for all  $X$ .
2. Well-separated interior configurations yield zero penalty.

### 3.5.3 Exercise 3: Symmetry

**Invariant:** permuting circle indices leaves  $E_{\text{pair}}$  unchanged.

### 3.5.4 Exercise 4: Proposal Locality

**Invariants:**

1. Exactly one circle moves per proposal.
2. Empirical variance of steps scales with  $\sigma^2$ .

### 3.5.5 Exercise 5: Metropolis Correctness

**Invariants:**

1. Downhill moves are always accepted.
2. Uphill acceptance frequencies match  $e^{-\Delta E/T}$ .

### 3.5.6 Exercise 6: Best-So-Far Monotonicity

**Invariant:** the best recorded energy is nonincreasing over time.

### 3.5.7 Exercise 7: Temperature Effects

**Observation:** higher initial temperature improves escape from overlaps on small  $N$ .

## 3.6 Diagnostics and Instrumentation

Acceptance rate is a primary diagnostic:

- early phase: 0.2–0.6,
- late phase: 0.01–0.2.

Deviations indicate mismatches between proposal scale, temperature, and energy magnitude.

## 3.7 Minimal Interfaces

The following interfaces are intentionally narrow:

```
energy.energy(X, L, r, alpha) -> float
energy.pair_energy(X, r) -> float
energy.wall_energy(X, L, r) -> float

propose.propose_move(X, L, r, sigma, rng)
-> (X_new, moved_index)

metropolis.accept(delta_E, T, rng) -> bool

anneal.run(X0, L0, r, alpha, T0, gamma, n_steps, sigma, seed)
-> dict(trace, best_state, best_energy, accept_rate)
```

Each function corresponds to a single theoretical concept from Chapter 1.

## 3.8 Implementation Notes

- Keep functions pure; pass RNG objects explicitly.
- Prefer clarity over vectorization in the baseline.
- Log energy, acceptance, and temperature at every step.
- Add performance optimizations *only after* all tests pass.

### 3.9 Stretch Goals

1. Adaptive proposal scaling based on acceptance rate.
2. Two-stage annealing schedules.
3. Feasibility search via outer bisection in  $L$ .
4. Incremental energy updates for  $\mathcal{O}(N)$  proposals.

### 3.10 Checkpoint

At this stage, you should have a solver where:

- all tests pass deterministically;
- the algorithm reflects the assumptions of Chapter 1;
- failures are explainable via diagnostics;
- the code is ready to be translated into C in subsequent chapters.

This completes the transition from *theory* to *validated implementation*. In Chapter 3, we repeat this process for gradient-based optimization, culminating in a fully test-driven C implementation of Adam.

## Chapter 4

# Inquiry-Based Implementation of Stochastic Gradient Descent with Adam

### Guiding Question

How can noisy, partial gradient information be systematically transformed into a stable, scalable optimization method for high-dimensional, nonconvex problems?

This chapter develops stochastic gradient descent (SGD) and the Adam optimizer through an inquiry-based sequence. The emphasis is on understanding the mathematical role of noise, momentum, and adaptivity, and on implementing Adam correctly from first principles.

### 4.1 From Deterministic to Stochastic Gradients

Consider the finite-sum optimization problem

$$\min_{\theta \in \mathbb{R}^d} F(\theta) := \frac{1}{N} \sum_{i=1}^N f_i(\theta).$$

#### 4.1.1 Inquiry

- What happens if  $N$  is very large?
- Is it necessary to evaluate all  $f_i$  at every iteration?
- What if we replace the full gradient with an estimator?

This leads to stochastic gradients. Given a random index  $i_k$ , define

$$g_k := \nabla f_{i_k}(\theta_k).$$

Then  $\mathbb{E}[g_k] = \nabla F(\theta_k)$ , but  $g_k$  has nonzero variance.

## 4.2 Stochastic Gradient Descent

The SGD iteration is

$$\theta_{k+1} = \theta_k - \eta_k g_k.$$

### 4.2.1 Inquiry

- Why does SGD converge despite noisy gradients?
- What role does the learning rate  $\eta_k$  play?
- Why must  $\eta_k \rightarrow 0$  in theory but not always in practice?

### 4.2.2 Key Phenomena

- Gradient noise acts as implicit regularization.
- SGD escapes shallow local minima and saddle points.
- Variance limits the achievable accuracy.

## 4.3 Momentum as Time Averaging

To reduce variance and accelerate convergence, introduce momentum:

$$\begin{aligned} v_{k+1} &= \beta v_k + (1 - \beta) g_k, \\ \theta_{k+1} &= \theta_k - \eta v_{k+1}. \end{aligned}$$

### 4.3.1 Inquiry

- Why does averaging gradients reduce noise?
- How is momentum related to low-pass filtering?
- Why can momentum overshoot minima?

Momentum can be interpreted as discretizing a second-order differential equation with damping.

## 4.4 Adaptive Learning Rates

Different coordinates may have gradients of very different scales. Adaptive methods address this by normalizing updates.

Define the second-moment accumulator

$$s_{k+1} = \beta_2 s_k + (1 - \beta_2) g_k^2,$$

where the square is taken elementwise.

### 4.4.1 Inquiry

- Why does dividing by  $\sqrt{s_k}$  stabilize training?
- What happens when gradients are sparse?
- Why is per-coordinate adaptivity dangerous?

## 4.5 The Adam Optimizer

Adam combines momentum and adaptive scaling.

### 4.5.1 Algorithm

Initialize  $m_0 = 0$ ,  $v_0 = 0$ .

$$\begin{aligned} m_{k+1} &= \beta_1 m_k + (1 - \beta_1) g_k, \\ v_{k+1} &= \beta_2 v_k + (1 - \beta_2) g_k^2. \end{aligned}$$

Bias correction:

$$\hat{m}_{k+1} = \frac{m_{k+1}}{1 - \beta_1^{k+1}}, \quad \hat{v}_{k+1} = \frac{v_{k+1}}{1 - \beta_2^{k+1}}.$$

Update:

$$\theta_{k+1} = \theta_k - \eta \frac{\hat{m}_{k+1}}{\sqrt{\hat{v}_{k+1}} + \varepsilon}.$$

### 4.5.2 Inquiry

- Why is bias correction necessary?
- What goes wrong without  $\varepsilon$ ?
- Why does Adam often converge faster but generalize worse?

## 4.6 Test-Driven Implementation Strategy

We implement Adam using TDD to avoid silent bugs.

### 4.6.1 Core Invariants to Test

- Determinism given a fixed random seed.
- Shape consistency of all tensors.
- Non-negativity of second-moment estimates.
- Correct bias correction at small  $k$ .

## 4.7 Exercise Sequence (TDD)

### 4.7.1 Exercise 1: Gradient Oracle

Write a test that checks a stochastic gradient estimator is unbiased on a quadratic function.

### 4.7.2 Exercise 2: Momentum Averaging

Verify that  $m_k$  equals an exponential moving average of past gradients.

### 4.7.3 Exercise 3: Second-Moment Accumulator

Test that  $v_k$  tracks the empirical variance scale of gradients.

### 4.7.4 Exercise 4: Bias Correction

On a constant gradient, verify that  $\hat{m}_k$  converges immediately to the true gradient.

### 4.7.5 Exercise 5: Adam Step on Quadratic Bowl

Check that Adam converges to the minimizer of  $f(\theta) = |\theta|^2$  from random initialization.

## 4.8 Failure Modes

Symptom	Cause
extbfSymptom	<b>Cause</b>
Divergence	Learning rate too large
height	$\beta_1$ too small
Slow convergence	$\beta_2$ too large
Parameter drift	
Poor generalization	Excessive adaptivity

extbf{Symptom}	<b>Cause</b>
Divergence	Learning rate too large
Slow convergence	$\beta_1$ too small
Parameter drift	$\beta_2$ too large
Poor generalization	Excessive adaptivity

# Chapter 5

## Test-Driven Adam in C (From Scratch)

### Purpose

This chapter is a correctness-first, test-driven implementation of the Adam optimizer in C. The emphasis is on:

- a minimal, auditable implementation (no external ML frameworks),
- deterministic behavior and numerical stability,
- unit tests that catch the common silent bugs (bias correction, moments, epsilon placement, shape/stride errors),
- extensibility toward HPC settings (SIMD/OpenMP) *after* tests are green.

### Guiding Question

How do we implement Adam in plain C so that (i) every mathematical step is testable, and (ii) the code is suitable as a performance baseline for later optimization?

#### 5.1 Scope and Assumptions

We optimize parameters  $\theta \in \mathbb{R}^d$  given a gradient vector  $g \in \mathbb{R}^d$  supplied by a *gradient oracle*.

This chapter focuses on the optimizer only. We will implement:

- Adam state and update step;
- optional weight decay (decoupled AdamW form);

- float/double support via a typedef;
- a tiny test harness (no third-party dependencies required).

## 5.2 Mathematical Specification

Given hyperparameters  $\eta > 0$ ,  $\beta_1, \beta_2 \in (0, 1)$ , and  $\varepsilon > 0$ , Adam performs:

$$\begin{aligned} m_t &= \beta_1 m_{t-1} + (1 - \beta_1) g_t, \\ v_t &= \beta_2 v_{t-1} + (1 - \beta_2) g_t^{\odot 2}, \\ \hat{m}_t &= \frac{m_t}{1 - \beta_1^t}, \quad \hat{v}_t = \frac{v_t}{1 - \beta_2^t}, \\ \theta_{t-1} &\leftarrow \theta_{t-1} - \eta \frac{\hat{m}_t}{\sqrt{\hat{v}_t} + \varepsilon}. \end{aligned}$$

All operations are elementwise except scalar multiplications.

### 5.2.1 AdamW (Decoupled Weight Decay)

If using AdamW with weight decay  $\lambda \geq 0$ :

$$\theta_{t-1} \leftarrow \theta_{t-1} - \eta \lambda \theta_{t-1}$$

in addition to the Adam step above. (Decoupled decay is preferred over adding  $\lambda\theta$  to gradients.)

## 5.3 Design Requirements (What Must Be True)

1. **Bias correction correctness:** first few steps must match the analytic formulas.
2. **Second moment non-negativity:**  $v_t[i] \geq 0$  always.
3. **Epsilon placement:**  $\varepsilon$  is added *outside* the square root:  $\sqrt{\hat{v}} + \varepsilon$ .
4. **No hidden allocations:** all buffers are allocated by the caller or in init.
5. **Deterministic stepping:** given the same inputs, the update is bitwise deterministic (within floating-point expectations across compilers).

## 5.4 Project Skeleton

A minimal repository layout:

```
adam_c/
include/
adam.h
adam_types.h
testHarness.h
src/
adam.c
adam_init.c
tests/
test_adam_bias.c
test_adam_moments.c
test_adam_quadratic.c
test_adam_adamw.c
Makefile
```

## 5.5 C API (Small and Testable)

### 5.5.1 Type Configuration

Use a single scalar type.

```
// adam_types.h
#pragma once
#include <stddef.h>

#ifndef ADAM_SCALAR_T
#define ADAM_SCALAR_T double
#endif

typedef ADAM_SCALAR_T adam_t;
```

### 5.5.2 Optimizer State

```
// adam.h
#pragma once
#include "adam_types.h"

typedef struct {
    size_t d;           // dimension
    adam_t lr;          // eta
```

```

adam_t beta1;
adam_t beta2;
adam_t eps;
adam_t weight_decay; // lambda (AdamW); 0 disables

// time step (starts at 0, increments on each step)
unsigned long long t;

// moments
adam_t *m; // length d
adam_t *v; // length d

// cached powers for bias correction
adam_t beta1_pow; // beta1^t
adam_t beta2_pow; // beta2^t
} adam_opt_t;

// init: caller provides buffers m and v of length d
int adam_init(adam_opt_t *opt, size_t d,
               adam_t lr, adam_t beta1, adam_t beta2, adam_t eps,
               adam_t weight_decay,
               adam_t *m_buf, adam_t *v_buf);

// reset moments and time
void adam_reset(adam_opt_t *opt);

// single step: t\theta and grad are length d
void adam_step(adam_opt_t *opt, adam_t *t\theta, const adam_t *grad);

```

## 5.6 Implementation (Reference)

`src/adaminit.c`

```

#include "adam.h"

static void zero_vec(adam_t *x, size_t d) {
    for (size_t i = 0; i < d; ++i) x[i] = (adam_t)0;
}

```

```

int adam_init(adam_opt_t *opt, size_t d,
adam_t lr, adam_t beta1, adam_t beta2, adam_t eps,
adam_t weight_decay,
adam_t *m_buf, adam_t *v_buf) {
if (!opt || !m_buf || !v_buf || d == 0) return -1;
opt->d = d;
opt->lr = lr;
opt->beta1 = beta1;
opt->beta2 = beta2;
opt->eps = eps;
opt->weight_decay = weight_decay;
opt->t = 0;
opt->m = m_buf;
opt->v = v_buf;
opt->beta1_pow = (adam_t)1;
opt->beta2_pow = (adam_t)1;
zero_vec(opt->m, d);
zero_vec(opt->v, d);
return 0;
}

```

```

void adam_reset(adam_opt_t *opt) {
if (!opt) return;
opt->t = 0;
opt->beta1_pow = (adam_t)1;
opt->beta2_pow = (adam_t)1;
for (size_t i = 0; i < opt->d; ++i) {
opt->m[i] = (adam_t)0;
opt->v[i] = (adam_t)0;
}
}

```

src/adam.c

```

#include "adam.h"
#include <math.h>

void adam_step(adam_opt_t *opt, adam_t *theta, const adam_t *grad) {
const size_t d = opt->d;

```

```
// t := t + 1 and update cached powers
opt->t += 1;
opt->beta1_pow *= opt->beta1;
opt->beta2_pow *= opt->beta2;

const adam_t one = (adam_t)1;
const adam_t b1 = opt->beta1;
const adam_t b2 = opt->beta2;
const adam_t lr = opt->lr;
const adam_t eps = opt->eps;

const adam_t inv_bias1 = one / (one - opt->beta1_pow);
const adam_t inv_bias2 = one / (one - opt->beta2_pow);

// Optional decoupled weight decay (AdamW)
if (opt->weight_decay != (adam_t)0) {
    const adam_t wd = opt->weight_decay;
    for (size_t i = 0; i < d; ++i) {
        t\theta[i] -= lr * wd * t\theta[i];
    }
}

// Moment updates + parameter step
for (size_t i = 0; i < d; ++i) {
    const adam_t g = grad[i];

    // m_t, v_t
    const adam_t m = b1 * opt->m[i] + (one - b1) * g;
    const adam_t v = b2 * opt->v[i] + (one - b2) * (g * g);
    opt->m[i] = m;
    opt->v[i] = v;

    // bias corrected
    const adam_t mhat = m * inv_bias1;
    const adam_t vhat = v * inv_bias2;

    // epsilon outside sqrt
```

```
t\theta[i] -= lr * (mhat / (sqrt(vhat) + eps));  
}  
}
```

## 5.7 A Tiny Test Harness (No Dependencies)

*include/test\_harness.h*

```
#pragma once  
  
#include <math.h>  
#include <stdio.h>  
#include <stdlib.h>  
  
  
#define ASSERT_TRUE(cond) do {  
    if (!(cond)) {  
        fprintf(stderr, "ASSERT_TRUE failed: %s (%s:%d)\n", #cond, **FILE**, **LINE**);  
        exit(1);  
    }  
} while (0)  
  
  
#define ASSERT_NEAR(a,b,tol) do {  
    double _da = (double)(a);  
    double _db = (double)(b);  
    double _dt = fabs(_da - _db);  
    if (_dt > (tol)) {  
        fprintf(stderr, "ASSERT_NEAR failed: |%s-%s|=%g > %g (%s:%d)\n", #a, #b, _dt, (double)(tol),  
        exit(1);  
    }  
} while (0)
```

## 5.8 TDD Exercise Sequence

### 5.8.1 Exercise 1: Bias Correction on Constant Gradient

**Goal:** catch the most common Adam bug.

Set  $g_t \equiv g$  constant,  $m_0 = v_0 = 0$ . Show analytically:

$$m_t = (1 - \beta_1^t)g, \quad \hat{m}_t = g.$$

Similarly  $\hat{v}_t = g^2$ . Therefore the update should be

$$\theta_t = \theta_{t-1} - \eta, \frac{g}{|g| + \varepsilon}.$$

**Test:** in 1D, with  $g = 2$ , verify the very first step matches the formula within tolerance.

**tests/test\_adambias.c**

```
#include "adam.h"
#include "test_harness.h"

int main(void) {
    adam_t m[1], v[1];
    adam_opt_t opt;
    ASSERT_TRUE(adam_init(&opt, 1, (adam_t)0.1, (adam_t)0.9, (adam_t)0.999,
        (adam_t)1e-8, (adam_t)0.0, m, v) == 0);

    adam_t t\theta[1] = {(adam_t)1.0};
    const adam_t grad[1] = {(adam_t)2.0};

    adam_step(&opt, t\theta, grad);

    // For constant g, bias-corrected mhat=g and vhat=g^2 at t=1.
    const adam_t expected = (adam_t)1.0 - (adam_t)0.1 * ((adam_t)2.0 / ((adam_t)2.0 + (adam_t)1e-
        ASSERT_NEAR(t\theta[0], expected, 1e-12);

    return 0;
}
```

### 5.8.2 Exercise 2: Second Moment Non-Negativity

**Test:** feed arbitrary gradients and verify  $v[i] \geq 0$  always.

**tests/test\_adammoments.c**

```
#include "adam.h"
#include "test_harness.h"

int main(void) {
    enum { d = 8 };
    adam_t m[d], v[d];
```

```

adam_opt_t opt;
ASSERT_TRUE(adam_init(&opt, d, (adam_t)1e-2, (adam_t)0.9, (adam_t)0.99,
(adam_t)1e-8, (adam_t)0.0, m, v) == 0);

adam_t t\thetaeta[d];
adam_t g[d];
for (int i = 0; i < d; ++i) { t\thetaeta[i] = (adam_t)0; g[i] = (adam_t)(i - 3); }

for (int k = 0; k < 100; ++k) {
for (int i = 0; i < d; ++i) g[i] = (adam_t)((k + 1) * (i - 3));
adam_step(&opt, t\thetaeta, g);
for (int i = 0; i < d; ++i) ASSERT_TRUE(opt.v[i] >= (adam_t)0);
}

return 0;
}

```

### 5.8.3 Exercise 3: Epsilon Placement Regression Test

**Goal:** ensure  $\sqrt{\hat{v}} + \varepsilon$  and not  $\sqrt{\hat{v} + \varepsilon}$ .

**Test idea:** choose  $\hat{v}$  extremely small (e.g. gradient nearly zero) so that the two expressions differ measurably.

### 5.8.4 Exercise 4: AdamW Decoupled Weight Decay

**Test:** with  $g = 0$ , Adam step should be zero but AdamW should decay parameters by

$$\theta \leftarrow (1 - \eta\lambda)\theta.$$

`tests/testadamadamw.c`

```

#include "adam.h"
#include "test_harness.h"

int main(void) {
adam_t m[1], v[1];
adam_opt_t opt;
const adam_t lr = (adam_t)0.1;
const adam_t wd = (adam_t)0.5;
ASSERT_TRUE(adam_init(&opt, 1, lr, (adam_t)0.9, (adam_t)0.999,
(adam_t)1e-8, wd, m, v) == 0);

```

```

adam_t t\theta[1] = {(adam_t)2.0};
const adam_t grad[1] = {(adam_t)0.0};

adam_step(&opt, t\theta, grad);

const adam_t expected = (adam_t)2.0 * ((adam_t)1.0 - lr * wd);
ASSERT_NEAR(t\theta[0], expected, 1e-12);
return 0;
}

```

### 5.8.5 Exercise 5: End-to-End on a Quadratic Bowl

Minimize

$$f(\theta) = \frac{1}{2}|\theta|_2^2 \Rightarrow \nabla f(\theta) = \theta.$$

**Test:** initialize  $\theta_0$  and run many Adam steps with  $g_t = \theta_t$ ; verify  $|\theta|$  decreases below a threshold.

tests/test<sub>adam</sub>quadratic.c

```

#include "adam.h"
#include "test_harness.h"

static adam_t norm2(const adam_t *x, size_t d) {
    adam_t s = (adam_t)0;
    for (size_t i = 0; i < d; ++i) s += x[i]*x[i];
    return (adam_t)sqrt((double)s);
}

int main(void) {
    enum { d = 4 };
    adam_t m[d], v[d];
    adam_opt_t opt;
    ASSERT_TRUE(adam_init(&opt, d, (adam_t)1e-1, (adam_t)0.9, (adam_t)0.999,
        (adam_t)1e-8, (adam_t)0.0, m, v) == 0);

    adam_t t\theta[d] = {(adam_t)5.0, (adam_t)-3.0, (adam_t)2.0, (adam_t)-1.0};
    adam_t g[d];

```

```

const adam_t n0 = norm2(t\theta, d);
for (int k = 0; k < 2000; ++k) {
    for (int i = 0; i < d; ++i) g[i] = t\theta[i];
    adam_step(&opt, t\theta, g);
}
const adam_t n1 = norm2(t\theta, d);
ASSERT_TRUE(n1 < (adam_t)1e-2 * n0);
return 0;
}

```

## 5.9 Makefile (Minimal)

```

CC ?= gcc
CFLAGS ?= -O2 -std=c11 -Wall -Wextra -Iinclude
LDFLAGS ?= -lm

SRC = src/adam.c src/adam_init.c

TESTS =
tests/test_adam_bias
tests/test_adam_moments
tests/test_adam_adamw
tests/test_adam_quadratic

all: $(TESTS)

tests/%: tests/%.c $(SRC)
$(CC) $(CFLAGS) -o $@ $^ $(LDFLAGS)

test: all
@for t in $(TESTS); do echo "[RUN] $$t"; $$t; echo "[OK ] $$t"; done

clean:
rm -f $(TESTS)

```

## 5.10 Debugging Checklist

When a test fails, the most likely issues are:

- using  $\beta^t$  with wrong  $t$  indexing (off-by-one);
- bias-correction computed using the *old* power rather than updated power;
- epsilon placed inside the square root;
- integer truncation when computing norms or tolerances;
  - accidental aliasing between exttt $\theta$  and extttgradbuffers.

## 5.11 HPC Extensions (After Tests are Green)

1. **Incremental vectorization:** replace the inner loop with SIMD intrinsics.
2. **OpenMP:** parallelize the parameter dimension for large  $d$ .
3. **Mixed precision:** keep moments in float, accumulate in double.
4. **Fused kernels:** combine weight decay, moments, and update in one pass.

## 5.12 Checkpoint

By the end of this chapter you should have a C repository where:

- extttmake test executes all tests and returns success;
- the Adam step matches analytic bias-corrected formulas at early iterations;
- the quadratic-bowl end-to-end test demonstrates stable convergence;
- the code is ready to be optimized without sacrificing correctness.

## Chapter 6

# Optimizing the Square Size: Feasibility and Joint Optimization

Up to this point, simulated annealing has been used to optimize circle positions inside a *fixed* square of side length  $L$ . We now address a more fundamental geometric question:

What is the smallest square side length  $L^*$  for which  $N$  identical circles of radius  $r$  can be packed without overlap?

This chapter develops two complementary solver strategies:

1. **Feasibility search:** treat  $L$  as a parameter, test feasibility via SA, and search for the minimal feasible  $L$ .
2. **Joint optimization:** include  $L$  directly in the optimization state and minimize a soft energy that trades off feasibility and size.

Both approaches are built using the same inquiry-driven workflow:

Geometry  $\longrightarrow$  Invariant  $\longrightarrow$  Test  $\longrightarrow$  Code.

Throughout, we distinguish between:

- *guaranteed bounds* derived from geometry, and
- *heuristic behavior* introduced by stochastic optimization.

### 6.1 Why Optimizing $L$ Is Hard

Packing problems are inherently nonconvex. Even for fixed  $L$ , feasibility is not a simple constraint satisfaction problem; it must be discovered algorithmically.

Optimizing  $L$  adds a second layer of difficulty:

- reducing  $L$  increases boundary pressure,
- smaller  $L$  amplifies overlap penalties,
- feasibility becomes a probabilistic outcome of the inner solver.

We therefore separate the problem into two solver paradigms.

## 6.2 Approach I: Feasibility Search

### 6.2.1 Problem Formulation

For a fixed  $L$ , define *feasibility* as:

$$E_{\text{pair}}(X) = 0 \quad \text{and} \quad E_{\text{wall}}(X, L) = 0.$$

Numerically, we test:

$$E_{\text{pair}}(X) \leq \varepsilon, \quad E_{\text{wall}}(X, L) \leq \varepsilon,$$

for a small tolerance  $\varepsilon$ .

The feasibility solver asks:

For which values of  $L$  does simulated annealing find a feasible configuration?

This motivates an outer search in  $L$ .

### 6.2.2 Guaranteed Lower Bounds

Before running any optimization, geometry alone already constrains  $L$ .

#### Diameter Bound

At minimum, a single circle must fit:

$$L \geq 2r.$$

This is a trivial but unavoidable constraint.

#### Area Bound

The total area of all circles is:

$$A_{\text{circles}} = N\pi r^2.$$

Since the square has area  $L^2$ , any feasible packing must satisfy:

$$L^2 \geq N\pi r^2 \Rightarrow L \geq r\sqrt{N\pi}.$$

This bound is *necessary* and therefore guaranteed.

### Density-Aware Bound (Optional)

The maximal packing density of equal circles in the plane is

$$\delta_{\max} = \frac{\pi}{2\sqrt{3}} \approx 0.9069.$$

Even ignoring boundary effects,

$$L^2 \geq \frac{N\pi r^2}{\delta_{\max}} \Rightarrow L \geq r\sqrt{\frac{N\pi}{\delta_{\max}}}.$$

Boundary effects only reduce achievable density, so this remains a valid lower bound.

**Combined lower bound.** In practice, we use:

$$L_{\text{lo}} = \max\left(2r, r\sqrt{N\pi}\right),$$

optionally replacing the area bound with the density-aware bound.

### 6.2.3 Guaranteed Upper Bounds

Lower bounds alone are insufficient; feasibility search also requires a value of  $L$  that is *known* to be feasible.

#### Grid Construction

A simple constructive packing places circles on a square grid with spacing  $2r$ :

$$x_{ij} = (r + 2ri, r + 2rj).$$

Let

$$k = \lceil \sqrt{N} \rceil.$$

Then a  $k \times k$  grid fits inside a square of side length:

$$L_{\text{grid}} = 2rk.$$

This configuration has:

$$E_{\text{pair}} = 0, \quad E_{\text{wall}} = 0,$$

and therefore provides a *constructive feasibility witness*.

**Upper bound.** We set:

$$L_{\text{hi}} = 2r \lceil \sqrt{N} \rceil.$$

This bound is crude but guaranteed and deterministic.

#### 6.2.4 Outer Bisection on $L$

With a bracket  $[L_{\text{lo}}, L_{\text{hi}}]$  in hand, we perform an outer bisection search.

1. Set  $L_{\text{mid}} = \frac{1}{2}(L_{\text{lo}} + L_{\text{hi}})$ .
2. Run simulated annealing at fixed  $L_{\text{mid}}$ .
3. If a feasible configuration is found:
  - record the witness,
  - set  $L_{\text{hi}} \leftarrow L_{\text{mid}}$ .
4. Otherwise, set  $L_{\text{lo}} \leftarrow L_{\text{mid}}$ .

Because simulated annealing is stochastic, feasibility is a *probabilistic predicate*. To reduce false negatives, the solver may retry each  $L_{\text{mid}}$  with multiple deterministic seeds.

The algorithm terminates when:

$$L_{\text{hi}} - L_{\text{lo}} \leq \text{tolerance}.$$

### 6.3 Approach II: Joint Optimization of $(X, L)$

An alternative is to optimize positions and box size simultaneously.

#### 6.3.1 Energy Model

We define a soft objective:

$$E(X, L) = E_{\text{pair}}(X) + E_{\text{wall}}(X, L) + \alpha L,$$

with  $\alpha > 0$  controlling the pressure to shrink the square.

Key properties:

- $E(X, L) \geq 0$  always,
- feasibility corresponds to  $E_{\text{pair}} = E_{\text{wall}} = 0$ ,
- reducing  $L$  is rewarded only when geometry allows it.

### 6.3.2 Joint Proposal Kernel

The state space is now  $(X, L)$ . At each SA step:

- with probability  $p$ , propose a move of a single circle position,
- with probability  $1 - p$ , propose a move in  $L$ :

$$L' = \max(2r, L + \delta_L), \quad \delta_L \sim \mathcal{N}(0, \sigma_L^2).$$

Acceptance is handled by the same Metropolis rule using the correct  $\Delta E$ .

### 6.3.3 Interpretation

Joint optimization replaces hard feasibility with a soft competition:

- early at high temperature,  $L$  can shrink aggressively,
- overlaps and wall violations are temporarily tolerated,
- at low temperature, only feasible shrinkage survives.

Unlike feasibility search, this approach produces a *single trajectory* rather than an outer–inner loop, but its outcome depends sensitively on  $\alpha$  and the proposal scales.

## 6.4 Comparison of the Two Approaches

	Feasibility Search	Joint Optimization
Guarantees	Explicit feasibility	Soft, parameter-dependent
Bounds	Required	Not required
Complexity	Outer $\times$ inner	Single SA run
Interpretability	Clear witnesses	Continuous trade-off
Sensitivity	Low (with retries)	High ( $\alpha, \sigma_L$ )

In practice:

- feasibility search is preferable when correctness and reproducibility matter;
- joint optimization is useful for exploratory or heuristic searches.

## 6.5 Checkpoint

At the end of this chapter, you should be able to:

- derive guaranteed lower and upper bounds for  $L^*$ ;
- construct feasibility witnesses without optimization;
- implement an outer bisection loop driven by simulated annealing;
- understand the trade-offs of joint  $(X, L)$  optimization;
- distinguish mathematical guarantees from heuristic behavior.

In the next chapter, we revisit these solvers from a gradient-based perspective and contrast stochastic annealing with deterministic optimization methods.

## Chapter 7

# C Implementation of Feasibility and Joint Optimization Solvers

The previous chapter developed two solver strategies for optimizing the square side length:

1. feasibility search via an outer loop on  $L$ ,
2. joint optimization of  $(X, L)$  using simulated annealing.

This chapter translates those ideas into a *minimal, test-driven C implementation* that extends the existing simulated annealing baseline.

As before, the emphasis is not performance, but:

- correctness,
- explicit invariants,
- deterministic behavior under fixed seeds,
- and instrumentation suitable for scientific reasoning.

All additions are layered *on top of* the existing SA core, rather than rewriting it.

### 7.1 Design Principles for the Extension

Before writing code, we make several design commitments.

1. **No changes to the inner SA logic unless required.** The fixed- $L$  simulated annealing kernel remains valid and reusable.
2. **Feasibility is a predicate, not a heuristic.** It is expressed as a function of  $(X, L)$  and tested independently.

3. **Bounds come from geometry, not tuning.** Lower and upper bounds for  $L$  are computed deterministically.
4. **Solvers compose existing components.** New solvers call the existing SA driver instead of duplicating logic.

These principles keep the codebase modular and auditable.

## 7.2 New Supporting Primitives

### 7.2.1 Feasibility Predicate

The most basic new operation is a feasibility test.

**Definition.** A configuration  $(X, L)$  is feasible if:

$$E_{\text{pair}}(X) = 0 \quad \text{and} \quad E_{\text{wall}}(X, L) = 0,$$

up to a small numerical tolerance  $\varepsilon$ .

#### C interface.

```
int is_feasible(const Vec2* X, size_t N,
                double r, double L, double eps);
```

#### Design notes.

- This function is pure and deterministic.
- It relies only on energy components already implemented.
- It is used by:
  - the feasibility solver,
  - tests for geometric sanity,
  - post-run validation.

#### Tests.

- Known non-overlapping interior configurations return feasible.
- Any overlap or wall violation returns infeasible.
- Degenerate cases ( $N = 1, L < 2r$ ) behave correctly.

### 7.2.2 Lower and Upper Bounds on $L$

To initialize feasibility search, we compute a deterministic bracket.

#### Lower bounds

Two guaranteed lower bounds are implemented:

- Diameter bound:

$$L \geq 2r.$$

- Area bound:

$$L \geq r\sqrt{N\pi}.$$

In code:

```
double lower_bound_area(size_t N, double r);
double lower_bound_basic(size_t N, double r);
```

The solver uses:

$$L_{\text{lo}} = \max(2r, r\sqrt{N\pi}).$$

#### Upper bound via grid construction

A constructive packing places circles on a square grid with spacing  $2r$ .

#### C interface.

```
double upper_bound_grid(size_t N, double r);
void init_grid(Vec2* X, size_t N, double r, double L);
```

This guarantees:

$$L_{\text{hi}} = 2r\lceil\sqrt{N}\rceil$$

and provides an explicit feasible witness.

#### Tests.

- Grid initialization yields zero pair and wall energy.
- Returned  $L_{\text{hi}} \geq L_{\text{lo}}$  for all  $N$ .

## 7.3 Feasibility Solver in C

### 7.3.1 High-Level Structure

The feasibility solver performs an outer bisection on  $L$ , using the existing fixed- $L$  SA kernel as a feasibility oracle.

#### Interface.

```
FeasibleSolveResult solve_feasible_bisect(
    size_t N, double r,
    double L_lo, double L_hi,
    double eps feas,
    double T0, double gamma,
    size_t n_steps, double sigma,
    unsigned long seed,
    int retries_per_L
);
```

#### Algorithm.

1. Initialize  $(L_{\text{lo}}, L_{\text{hi}})$  from geometric bounds.
2. While  $L_{\text{hi}} - L_{\text{lo}}$  exceeds tolerance:
  - (a) Set  $L_{\text{mid}} = \frac{1}{2}(L_{\text{lo}} + L_{\text{hi}})$ .
  - (b) Run SA at fixed  $L_{\text{mid}}$ .
  - (c) If a feasible configuration is found:
    - record the witness,
    - set  $L_{\text{hi}} \leftarrow L_{\text{mid}}$ .
  - (d) Otherwise set  $L_{\text{lo}} \leftarrow L_{\text{mid}}$ .

**Handling stochasticity.** Because SA is probabilistic, feasibility tests may fail spuriously. The solver supports multiple retries per  $L$ , using deterministic seed offsets to preserve reproducibility.

**Diagnostics.** Each bisection iteration may log:

- tested  $L$ ,
- feasibility result,
- best energy achieved,
- best  $E_{\text{pair}}$  and  $E_{\text{wall}}$ .

## 7.4 Joint Optimization Solver in C

### 7.4.1 Extended State Space

Joint optimization treats  $(X, L)$  as a single state.

**State variables.**

- $X \in \mathbb{R}^{N \times 2}$ ,
- $L \geq 2r$ .

**Energy.**

$$E(X, L) = E_{\text{pair}}(X) + E_{\text{wall}}(X, L) + \alpha L.$$

### 7.4.2 Proposal Kernel

At each step:

- with probability  $p$ , propose a circle move (as in the baseline),
- with probability  $1 - p$ , propose:

$$L' = \max(2r, L + \delta_L), \quad \delta_L \sim \mathcal{N}(0, \sigma_L^2).$$

Only the wall energy and  $\alpha L$  term change for an  $L$ -move.

### 7.4.3 C Interface

```
JointAnnealResult anneal_joint_run(
    const Vec2* X0, size_t N,
    double L0, double r, double alpha,
    double T0, double gamma,
    size_t n_steps,
    double sigma_x, double sigma_L,
    double p_move_x,
    unsigned long seed
);
```

**Trace logging.** The joint solver records:

- $E(t)$ ,

- $L(t)$ ,
- temperature,
- acceptance decisions,
- move type (position vs.  $L$ ).

This allows direct inspection of how  $L$  evolves during annealing.

## 7.5 Testing Strategy

The new solvers introduce both mathematical and algorithmic invariants.

### 7.5.1 Mathematical tests

- Lower bounds are never violated.
- Grid upper bound is always feasible.
- $\Delta E$  for  $L$ -moves matches recomputed energy differences.

### 7.5.2 Algorithmic tests

- Bisection interval shrinks monotonically.
- Best-so-far feasible  $L$  never increases.
- Joint solver respects  $L \geq 2r$  at all times.

### 7.5.3 Empirical sanity checks

- Larger  $\alpha$  produces smaller final  $L$  on average.
- Acceptance rate decreases as temperature cools.

These are treated as regression tests with tolerant thresholds, not strict invariants.

## 7.6 Checkpoint

At the end of this chapter, you should have:

- a principled feasibility solver built from geometry and SA,
- a joint  $(X, L)$  annealer with explicit proposal semantics,

- deterministic, testable C code extending the original baseline,
- diagnostics that explain solver behavior rather than obscure it.

The next chapter turns to *performance*: spatial acceleration, parallelism, and the cost of correctness.