

# HPC Simulated Annealing for Non-Convex Polygon Packing in a Square

Bracketing, Bisection Refinement, and Time-Limited Shrink Polish

Joel Amir D. Maldonado Tänori

Applied Mathematics PhD, University of Arizona

January 15, 2026

# Agenda

- ① Problem and constraints
- ② Geometry model: triangulation + SAT + broad-phase rejects
- ③ Energy function, feasibility, and SA moves
- ④ Outer optimizer: bracketing + bisection
- ⑤ Time-limited polish: adaptive shrink search
- ⑥ HPC reliability: determinism, checkpoints, SIGTERM flush
- ⑦ Results / best-solution visuals

# Problem Statement

**Goal.** Pack  $N$  identical copies of a fixed **non-convex** polygon  $P \subset \mathbb{R}^2$  into a square of side length  $L$ .

**Decision variables (per instance  $i = 1, \dots, N$ ):**

$$(c_x^{(i)}, c_y^{(i)}, \theta^{(i)}) \in \mathbb{R}^2 \times [0, 2\pi)$$

where  $(c_x^{(i)}, c_y^{(i)})$  is translation and  $\theta^{(i)}$  is rotation.

**Constraints.**

- **Non-overlap:** interiors of instances are disjoint.
- **Containment:** each polygon lies inside the square  $[-L/2, L/2]^2$ .

**Objective.** Minimize  $L$  (tightest feasible packing).

# A Provable Lower Bound (Area Bound)

Let  $\text{area}(P)$  be the polygon area. Then any feasible packing must satisfy:

$$L^2 \geq N \cdot \text{area}(P) \Rightarrow L \geq \sqrt{N \text{area}(P)}.$$

In code, we keep a strictly provably infeasible threshold:

$$L_{\text{area}} = \sqrt{N \text{area}(P)}, \quad L_{\text{area-infeas}} = (1 - \varepsilon) L_{\text{area}}.$$

**Interpretation.** If  $L \leq L_{\text{area-infeas}}$ , the instance is **provably infeasible** (no need to run SA).

# Geometry Model: Triangulation + SAT

We represent the non-convex polygon  $P$  by:

- A fixed vertex list  $\{v_k\}_{k=1}^{N_V}$  in local coordinates.
- A fixed triangulation  $\{\Delta_t\}_{t=1}^{N_T}$ , each triangle uses indices into the vertex list.

Each instance  $i$  produces world vertices:

$$w_k^{(i)} = R(\theta^{(i)}) v_k + \begin{bmatrix} c_x^{(i)} \\ c_y^{(i)} \end{bmatrix}.$$

**Collision test between instances  $(i, j)$ :**

- ① Broad-phase AABB overlap
- ② Broad-phase bounding-circle reject
- ③ Narrow-phase triangle-triangle SAT for all  $(t_a, t_b)$

# Broad-Phase Acceleration

**AABB reject.** If axis-aligned bounding boxes do not overlap, polygons cannot overlap.

**Bounding-circle reject.** Precompute base radius  $r$ :

$$r = \max_k \|v_k\|_2.$$

If centers are far:

$$\|c^{(i)} - c^{(j)}\|_2 > 2r \quad \Rightarrow \quad \text{no overlap.}$$

**Uniform grid (spatial hashing).** Each instance belongs to one grid cell; collision checks only consider neighbor cells within a radius based on  $2r$ .

**Outcome.** Pairwise overlap checks scale closer to *local neighborhood* interactions rather than  $O(N^2)$  in practice.

# Containment Penalty via AABB

Containment is enforced through an **outside penalty** computed from AABB vs. square:

$$\text{out}_i(L) = \sum_{\text{violations}} d^2$$

where  $d$  is how far the AABB exceeds  $\pm L/2$  in any direction.

**Benefit.** Cheap to compute, differentiable enough for SA, and effective when combined with increasing penalty weights in Phase B.

# Energy Function and Feasibility Metric

We separate the concept of **energy** (for SA acceptance) from **feasibility** (for outer logic).

**Totals:**

$$\text{ov}(L) = \sum_{i < j} \text{overlap\_penalty}(i, j), \quad \text{out}(L) = \sum_i \text{outside\_penalty}(i).$$

**Energy (within an SA run at fixed  $L$ ):**

$$\mathcal{E} = \lambda \text{ov} + \mu \text{out}$$

(with weights scheduled by phase).

**Feasibility metric:**

$$\text{feas} = \text{ov} + \text{out}.$$

We declare “feasible” if  $\text{feas} \leq \tau$  for a small tolerance  $\tau$  and  $L$  is not area-provably infeasible.

# Move Set and Incremental Updates

Each SA iteration chooses a random index  $k$  and applies one of:

- **Reinsert (small probability):** randomize  $(c_x, c_y, \theta)$  uniformly.
- **Local jitter:**  $(c_x, c_y) \leftarrow (c_x, c_y) + \Delta$  with  $\Delta \sim \text{Unif}([-s, s]^2)$ .
- **Rotation mix:** with probability  $p_{\text{rot}}$ ,  $\theta \leftarrow \theta + \Delta\theta$ .

**Incremental bookkeeping.** Only terms involving instance  $k$  need recomputation:

$$\text{ov} \leftarrow \text{ov} + (\text{ov}_k^{\text{new}} - \text{ov}_k^{\text{old}}), \quad \text{out} \leftarrow \text{out} + (\text{out}_k^{\text{new}} - \text{out}_k^{\text{old}}).$$

**Outcome.** Fast inner loop; suitable for HPC sweeps over many  $N$ .

# Two-Phase SA Schedule

We use two sequential phases at fixed  $L$ :

## Phase A (Explore).

- Higher temperature range  $T_{\text{start}} \rightarrow T_{\text{end}}$
- Larger step sizes
- Moderate penalties  $(\lambda, \mu)$
- Purpose: escape poor initializations and reduce gross overlaps/outside

## Phase B (Enforce).

- Lower temperatures
- Smaller step sizes
- **Ramping** penalties every  $K$  iterations:  $(\lambda, \mu) \leftarrow \min((\lambda, \mu) \cdot \rho, (\lambda_{\max}, \mu_{\max}))$
- Purpose: aggressively drive feas  $\rightarrow 0$

## SA Acceptance Rule

Given current energy  $\mathcal{E}$  and proposed energy  $\mathcal{E}'$ , accept with:

$$\Delta\mathcal{E} = \mathcal{E}' - \mathcal{E}.$$

$$\mathbb{P}(\text{accept}) = \begin{cases} 1, & \Delta\mathcal{E} \leq 0, \\ \exp(-\Delta\mathcal{E}/T), & \Delta\mathcal{E} > 0. \end{cases}$$

**Key detail.** We track the **best feasibility** configuration seen in the run:

$$\text{feas}_{\min} = \min_t \text{feas}(t),$$

and restore to that best configuration at the end of the trial.

# Pseudocode: SA Trial at Fixed $L$

linewidth  
**SA\_Trial( $L$ )**

Initialize  
state  
 $x$   
(po-  
si-  
tions  
+  
an-  
gles),  
com-  
pute  
to

# Outer Loop Overview

We want the smallest feasible  $L$ . The outer optimizer proceeds as:

- ① **Initialize** a conservative  $L$  via grid layout (near-square arrangement).
- ② **Bracketing:**
  - If feasible at current  $L$ , shrink until infeasible to find  $[L_{\text{low}}, L_{\text{high}}]$ .
  - If infeasible, grow until feasible.
- ③ **Bisection:** refine the bracket to a tight feasible  $L$ .
- ④ **Polish:** time-limited stochastic descent shrinking  $L$  further.

**Reliability constraints.** Area bound prevents invalid brackets; best feasible configuration is carried across  $L$  updates via scaling.

# Bracketing Logic (with Area Bound)

We maintain:

$$L_{\text{low}} \text{ (infeasible), } L_{\text{high}} \text{ (feasible).}$$

**Case 1: initial feasible.** Repeatedly try  $L \leftarrow \alpha L$  with  $\alpha < 1$ :

- Stop if  $L \leq L_{\text{area-infeas}}$  (provably infeasible).
- Otherwise run SA at new  $L$ .
- If infeasible, set  $L_{\text{low}} \leftarrow L$ .

**Case 2: initial infeasible.** Repeatedly try  $L \leftarrow \beta L$  with  $\beta > 1$  until feasible; set  $L_{\text{high}}$ .

**Important detail.** We do *not* overwrite an SA-found infeasible lower bracket with the weaker area bound; we keep the tighter information.

# Bisection Refinement

Given a valid bracket  $L_{\text{low}} < L_{\text{high}}$ :

$$L_{\text{mid}} = \frac{1}{2}(L_{\text{low}} + L_{\text{high}}).$$

**If**  $L_{\text{mid}} \leq L_{\text{area-infeas}}$ : mark infeasible and set  $L_{\text{low}} \leftarrow L_{\text{mid}}$  without SA.

**Else:**

- Warm-start by scaling best-feasible configuration from  $L_{\text{high}}$  to  $L_{\text{mid}}$ .
- Run a bounded number of SA trials at  $L_{\text{mid}}$ .
- If feasible:  $L_{\text{high}} \leftarrow L_{\text{mid}}$  (and update best-feasible snapshot).
- Else:  $L_{\text{low}} \leftarrow L_{\text{mid}}$  and restore best-feasible state.

After  $k$  steps, bracket width shrinks by  $2^{-k}$ .

# Pseudocode: Bracket + Bisection

linewidth  
**Minimize**  
 $L$   
**(Bracket**  
+  
**Bi-**  
**sec-**  
**tion)**

Initialize  
 $L$   
from  
grid  
lay-  
out:

## Polish Stage: Adaptive Shrink Search

After bisection, we have a strong feasible configuration at  $L^*$ .

**Idea.** Repeatedly attempt a small shrink:

$$L_{\text{try}} = L^*(1 - \epsilon),$$

warm-start by scaling positions, then run a few SA trials.

**If feasible:** accept and update best solution.

**If infeasible:** reject and adjust  $\epsilon$ .

**Adaptive control.** Maintain a sliding window of attempts and tune  $\epsilon$  toward a target success rate (stochastic descent with backoff).

# Time Limit and Checkpoints

For HPC sweeps, polish is time-limited (e.g., 900s = 15 minutes) to guarantee completion.

## Periodic checkpoints (every $\Delta t$ seconds):

- Write csv/<prefix>.checkpoint\_Nxxx.csv
- Write img/<prefix>.checkpoint\_Nxxx.svg

## SIGTERM handling. If Slurm sends SIGTERM near walltime:

- Flush **best snapshot** to both best and checkpoint files.
- Exit cleanly.

## Outcome. Even canceled jobs yield usable artifacts.

# HPC Engineering: Determinism and File Safety

**Deterministic seeds.** Each trial seed derived from:

$$\text{seed} = f(\text{base\_seed}, \text{run\_id}, \text{trial\_id})$$

(using SplitMix64 diffusion), ensuring reproducibility across arrays.

**Unique output prefixes.** For SLURM arrays:

$$\text{out\_prefix} = N\{N\}_\text{job}\{\text{job}\}_\text{task}\{\text{task}\}$$

Prevents file collisions when tasks run concurrently.

**Compile-on-node option.** Avoid GLIBC mismatch by building on compute nodes (cluster dependent).

# Complexity Notes and Practical Performance

**Theoretical worst case.** Collision evaluation can be heavy, but practical performance is improved by:

- Uniform grid neighbor enumeration (local pairs)
- AABB + bounding-circle rejects (broad phase)
- Small fixed triangulation size ( $N_T$  constant)
- Incremental energy updates per move (only one instance changes)

**Scaling expectation.** Runtime grows with  $N$  mainly due to increased local density and neighbor interactions, not purely  $N^2$ .

# Result Artifacts Produced Per Task

For each  $N$ , the code writes:

- img/<prefix>\_best\_Nxxx.svg
- csv/<prefix>\_best\_polys\_Nxxx.csv
- img/<prefix>\_checkpoint\_Nxxx.svg
- csv/<prefix>\_checkpoint\_Nxxx.csv

The CSV header includes:

- prefix, run\_id, seed
- $L$  and best feasibility score
- $N$

**Recommendation.** For figures, use SVG exports directly in LaTeX (or convert to PDF for best typography).

## Example Best Solutions (Insert Your Figures)

[img/N10\\_jobXXXX\\_task10\\_best\\_N010.pdf](#)

[img/N50\\_jobXXXX\\_task50\\_best\\_N050.pdf](#)

# Summary

- Non-convex polygon packing solved with **two-phase SA** + robust geometry checks.
- **Area bound** provides provable infeasibility cutoff.
- **Bracket + bisection** yields a tight feasible  $L$  reliably.
- **Time-limited polish** improves  $L$  further under strict HPC budgets.
- Engineering choices ensure **reproducibility** and **artifact safety** on Slurm arrays.

## Next steps:

- Aggregate sweep results into an  $L^*(N)$  curve
- Compare against tiling-inspired warm-starts
- Increase SA trials for hard  $N$  and perform longer polish on selected cases

# Questions?